

Digital Image Processing



Lecture 04

Local Enhancement via Spatial Filtering

Dr. Muhammad Sajjad
R.A: Asad Ullah

Overview

7 Sharpening Spatial Filters

6 Smoothing Spatial Filters

Introduction to
Spatial
Filtering

1

Spatial
Filtering
(Convolution)

2

Example of
Convolution
Operation

3

Padding in
Spatial
Filtering

5 Convolution vs
Correlation


4

2

Introduction to Spatial Filtering

Spatial Filtering


- A technique used in image processing
- Each pixel's value is modified based on the pixel itself and the values of its neighboring pixels.
- Used for image enhancement, noise reduction, and edge detection.



Median Filter

Filter (or Kernel, Mask, Template, Window):


- The kernel is an array whose size defines the neighborhood of operation, and whose coefficients determine the nature of the filter.
- We prefer odd-sized kernels (e.g., 3x3, 5x5, etc.).
- Kernel size can be represented as $(m \times n)$, where m is number of rows of the kernel/filter and n is the number of columns.



Noisy image

Result of
Median
Filter

Cont.



$W (-1,-1)$	$W (-1,0)$	$W (-1,1)$
$W (0,-1)$	$W (0,0)$	$W (0,1)$
$W (1,-1)$	$W (1,0)$	$W (1,1)$

Filter Kernel, $w(s, t)$

Kernel coefficients


$f (x-1, y-1)$	$f (x-1, y)$	$f (x-1, y+1)$
$f (x, y-1)$	$f (x, y)$	$f (x, y+1)$
$f (x+1, y-1)$	$f (x+1, y)$	$f (x+1, y+1)$

Pixels values
under kernel
when it is
centered on (x, y)

Spatial Filtering (Convolution)

Linear Spatial Filtering

- Move the kernel across the image, pixel by pixel.
 - Perform a sum-of-products operation between the local region of the image f and the filter kernel K .
 - Place the computed value into the corresponding pixel in the output image.
 - Simple and fast.
 - Can blur the image and lose details
 - Examples: Smoothing(Box filter, Gaussian filter), Edge Detection (Sobel Filter)




Note: In convolution we rotate the kernel by 180 degree before product operation

Cont.

Non-Linear Filtering:


- Use operations other than the sum-of-products
- The output is based on operations that don't simply add up the pixel values.
- Example: selecting the median value in a neighborhood (e.g., Median filter) or selecting Min or Max value i.e. Max/Min filter.
- Also called Order-Statistic Filters
- Better at preserving details and removing noise such as salt-and-pepper noise.
- Slower and more complex.



Example of Convolution Operation


$$(f * w)(1, 1) = (2 * 0) + (3 * 1) + (5 * 0) + (1 * 1) + (0 * 0) + (2 * -1) + (3 * 0) + (4 * 1) + (2 * 0)$$

$$(f * w)(1, 1) = 6$$




Input Image

*



Kernel

Sum =



Output Image

Note: Here
the kernel is
pre-rotated by
180 degree

Example of Convolution Operation

$$(f * w)(1, 2) = (3 * 0) + (5 * 1) + (0 * 0) + (0 * 1) + (2 * 0) + (3 * -1) + (4 * 0) + (2 * 1) + (4 * 0)$$

$$(f * w)(1, 2) = 4$$

2	3	5	0	4
1	0	2	3	0
3	4	2	4	1
1	3	3	2	4
4	2	1	4	3

*

0	1	0
1	0	.1
0	1	0

Sum =

Input Image

Kernel

Output Image

Example of Convolution Operation

$$(f * w)(1, 3) = (5 * 0) + (0 * 1) + (4 * 0) + (2 * 1) + (3 * 0) + (0 * -1) + (2 * 0) + (4 * 1) + (1 * 0)$$

$$(f * w)(1, 3) = 6$$

2	3	5	0	4
1	0	2	3	0
3	4	2	4	1
1	3	3	2	4
4	2	1	4	3

*

0	1	0
1	0	-1
0	1	0

Sum =

	6	4	6	

Input Image

Kernel

Output Image

Example of Convolution Operation

$$(f * w)(2, 1) = (1 * 0) + (0 * 1) + (2 * 0) + (3 * 1) + (4 * 0) + (2 * -1) + (1 * 0) + (3 * 1) + (3 * 0)$$

$$(f * w)(2, 1) = 4$$

2	3	5	0	4
1	0	2	3	0
3	4	2	4	1
3	3	3	2	4
1	3	1	4	3

*

0	1	0
1	0	.1
0	1	0

Sum =

6	4	6	
4			

Input Image

Kernel

Output Image

Example of Convolution Operation

$$(f * w)(2, 2) = (0 * 0) + (2 * 1) + (3 * 0) + (4 * 1) + (2 * 0) + (4 * -1) + (3 * 0) + (3 * 1) + (2 * 0)$$

$$(f * w)(2, 2) = 5$$

2	3	5	0	4
1	0	2	3	0
3	4	2	4	1
1	3	3	2	4
4	2	1	4	3

*

0	1	0
1	0	.1
0	1	0

Sum =

6	4	6		
4	5			

Input Image

Kernel

Output Image

Example of Convolution Operation

$$(f * w)(2, 3) = (2 * 0) + (3 * 1) + (0 * 0) + (2 * 1) + (4 * 0) + (1 * -1) + (3 * 0) + (2 * 1) + (4 * 0)$$

$$(f * w)(2, 3) = 6$$

2	3	5	0	4
1	0	2	3	0
3	4	2	4	1
1	3	3	2	4
4	2	1	4	3

*

0	1	0
1	0	-1
0	1	0

Sum =

6	4	6	
4	5	6	

Input Image

Kernel

Output Image

Example of Convolution Operation

$$(f * w)(3, 1) = (3 * 0) + (4 * 1) + (2 * 0) + (1 * 1) + (3 * 0) + (3 * -1) + (4 * 0) + (2 * 1) + (1 * 0)$$

$$(f * w)(3, 1) = 4$$

2	3	5	0	4
1	0	2	3	0
3	4	2	4	1
3	4	2	2	4
1	3	3	4	3
4	2	1	4	3

*

0	1	0
1	0	.1
0	1	0

Sum =

6	4	6
4	5	6
4		

Input Image

Kernel

Output Image

Example of Convolution Operation

$$(f * w)(3, 2) = (4 * 0) + (2 * 1) + (4 * 0) + (3 * 1) + (3 * 0) + (2 * -1) + (2 * 0) + (1 * 1) + (4 * 0)$$

$$(f * w)(3, 2) = 4$$

2	3	5	0	4
1	0	2	3	0
3	4	2	4	1
1	3	3	2	4
4	2	1	4	3

*

0	1	0
1	0	.1
0	1	0

Sum =

6	4	6	
6	5	6	
4	4		

Input Image

Kernel

Output Image

Example of Convolution Operation

$$(f * w)(3,3) = (2 * 0) + (4 * 1) + (1 * 0) + (3 * 1) + (2 * 0) + (4 * -1) + (1 * 0) + (4 * 1) + (3 * 0)$$

$$(f * w)(3,3) = 7$$

2	3	5	0	4
1	0	2	3	0
3	4	2	4	1
1	3	3	2	4
4	2	1	4	3

*

0	1	0
1	0	-1
0	1	0

Sum =

6	4	6	
6	5	6	
4	5	7	
4	4	7	

Input Image

Kernel

Output Image

Padding in Spatial Filtering

- What is padding?
- Why we use padding in Spatial Filtering?
- In the figure shown, part of the kernel 'w' lies outside the image 'f' when processing the edge pixels. As a result, the sum of the products is undefined in that area.
- In the figure shown, the size of the output decreased when we didn't process the edge pixels.
- Padding ensures that the filter can be applied evenly across all pixels.

1	0	-1
1	0	-1
1	0	-1

w

0	4	1	3	0
1	3	1	0	2
4	0	2	4	3
2	0	0	1	4
0	3	1	2	0

f (no padding)

0	4	1	3	0
1	3	1	0	2
4	0	2	4	3
2	0	0	1	4
0	3	1	2	0

1	0	-1
1	0	-1
1	0	-1



1	0	-1
4	-2	-6
3	-4	-4

Padding in Spatial Filtering

Zero Padding

- Zero padding adds zeros around the image border to maintain the same size after filtering.
- **Use cases:** Commonly used in deep learning to ensure output feature maps match input dimensions, especially in multi-layer networks.

Replicate Padding

- Replicate padding copies border pixel values as padding to preserve edge information.
- **Use cases:** Useful in smoothing filters (e.g., Gaussian blur) to avoid artificial edges from zero padding.

For a kernel of size $m \times n$:

- Pad with $(m-1)/2$, rows of zeros (top and bottom).
- Pad with $(n-1)/2$, columns of zeros (left and right).

1	2	3
4	5	6
7	8	9

f (no padding)

1	0	-1
1	0	-1
1	0	-1

w

Here kernel (w) size is 3×3 :
 $3-1/2 = 1$ row above and below
 $3-1/2 = 1$ col left and right

Zero Padding

0	0	0	0	0
0	1	2	3	0
0	4	5	6	0
0	7	8	9	0
0	0	0	0	0

1	1	2	3	3
1	1	2	3	3
4	4	5	6	6
7	7	8	9	9
7	7	8	9	9

Replicate Padding


Convolution vs Correlation

Correlation

- Moving the center of a kernel over an image, and computing the sum of products at each location

$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t)f(x + s, y + t)$$

$$a = (m - 1)/2 \quad b = (n - 1)/2$$



Cont.


Convolution

- Moving 180° rotated kernel over an image, and computing the sum of products at each location.

$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t)f(x - s, y - t)$$

$$a = (m - 1)/2 \quad b = (n - 1)/2$$

Here the formula is slightly different. The minus “ - ” sign rotates the kernel by 180 degree.






Smoothing Spatial Filters

Used for: Blurring and noise reduction.

Blurring helps to:

- Blurring is usually used in preprocessing steps.
- Remove small details and noise before object extraction.
- Bridge small gaps in lines or curves.



Cont.

Averaging Linear Filters


- Replaces each pixel with the average of its neighboring pixels.
 - Image $M \times N$, Filter $m \times n$

$$g(x, y) = \frac{\sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x+s, y+t)}{\sum_{s=-a}^a \sum_{t=-b}^b w(s, t)}$$

	1	1	1
*	1	1	1
	1	1	1

Box Filter

All coefficients are equal



Cont.

Example of box kernel

Here the dark edges show zero padding




Original Image



3 x 3 box
kernel result



11 x 11 box
kernel result



21 x 21 box
kernel result

Cont.

Gaussian Filter (Weighted average linear filter)


- The filter where nearby pixels have higher weight, and distant pixels have lower weight.
- Image $M \times N$, Filter $m \times n$

$$g(x, y) = \frac{\sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)}{\sum_{s=-a}^a \sum_{t=-b}^b w(s, t)}$$

- Mathematically, Gaussian filter is;

$$w(s, t) = G(s, t) = K e^{-\frac{s^2 + t^2}{2\sigma^2}}$$

Here, K is the scale (i.e., 1 in the Weighted Average, as shown in the upper right figure), and σ is the standard deviation.


$$\frac{1}{16} * \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

Weighted Average
Nearby pixels have higher weight, and distant pixels have lower weight

$$\frac{1}{4.8976} * \begin{array}{|c|c|c|} \hline 0.3679 & 0.6065 & 0.3679 \\ \hline 0.6065 & 1.0000 & 0.6065 \\ \hline 0.3679 & 0.6065 & 0.3679 \\ \hline \end{array}$$

$[6\sigma] \times [6\sigma]$ can be used for the size of Gaussian kernel

Cont.

Example of Weighted Average Filters



Original Image



Result of 21×21
Gaussian kernel
with $\sigma = 3.5$



Result of 43×43
Gaussian kernel
with $\sigma = 7$


Add 1 to
make it
odd

$$[6 \times 3.5] \times [6 \times 3.5] = 21 \times 21$$


$$[6 \times 7] \times [6 \times 7] = 42+1 \times 42+1$$

Cont.


Example Comparison of Box and Gaussian Filter



Original Image



Result of 71x 71
box kernel



Result of 151 x 151
Gaussian kernel with
 $\sigma = 25$


Cont.

Order-statistic (Nonlinear) Filters

- Nonlinear
- Based on ordering (ranking) the pixels contained in the filter mask
- Replacing the value of the center pixel with the value determined by the ranking result
- E.g., median filter, max filter, min filter


Median Filtering

- Assigns the mid value of all the gray levels in the mask to the center of mask
- Useful in removing impulse noise (also known as salt-and-pepper-noise).




Cont.

Median Filtering




X-ray image of a
circuit board

Corrupted by salt-
and-pepper noise



Result 3x3
averaging filer



Using 3x3
median filter

Sharpening Spatial Filters

Definition:

- Sharpening filters enhance the edges and fine details of an image, making objects and boundaries more distinct.
- They work by emphasizing the high-frequency components of the image, such as edges.

Foundation:

- Blurring/smoothing is performed by spatial averaging (equivalent to integration).
- Sharpening is performed by noting only the gray level changes in the image that is the **differentiation**.

Applications:

- Edge Detection: Used in image segmentation, object recognition, and feature extraction.
- Image Enhancement: Improves clarity and sharpness, making details more visible.
- Computer Vision: Sharpening helps in recognizing finer details in objects, such as texture and structure.



Sharpening via Spatial Differentiation


First-Order Derivatives (Gradient-based sharpening):

- Measures the rate of change of pixel intensity.
- A basic definition of the first-order derivative of a one-dimensional function $f(x)$ is the difference,


$$\frac{\partial f}{\partial x} = f(x+1) - f(x)$$



Step Edges: Abrupt change



Ramp Edges: Gradual change



Roof Edges: Peak change.

The first derivative must be:

1. Zero in areas of constant intensity, (along flat segments)
2. Nonzero at the onset of an intensity step or ramp.
3. Nonzero along intensity ramps.

$$\begin{array}{c|c|c|c|c} 1 & 4 & 0 & 5 \end{array} * \begin{array}{c|c} -1 & 1 \end{array} \rightarrow \begin{array}{c|c|c} 3 & -4 & 5 \end{array}$$



First-order derivative kernel used for edge detection

Cont.

Second order derivatives of digital functions

We define the second-order derivative of $f(x)$ as the difference,

$$\frac{\partial^2 f}{\partial x^2} = f(x+1) + f(x-1) - 2f(x)$$

f(x-1)	f(x)	f(x+1)
--------	------	--------

1	-2	1
---	----	---

The second derivative must be:

- Zero in areas of constant intensity.
- Nonzero at the onset and end of an intensity step or ramp.
- Zero along intensity ramps.



1	4	0	5	3	2
---	---	---	---	---	---

1	-2	1
---	----	---

-7	9	-7	1
----	---	----	---

Example of derivatives

- 1st derivative detect thick edges while 2nd derivative detect thin edges.
- 2nd derivative has much stronger response at gray-level step than 1st derivative.



Various situations encountered for derivatives

$$f' = \frac{\partial f}{\partial x} \quad f'' = \frac{\partial^2 f}{\partial x^2}$$

- Ramps or steps in the 1D profile normally characterize the edges in an image
- f'' is nonzero at the onset and end of the ramp: produce thin (double) edges
- f' is nonzero along the entire ramp produce thick edges


• Flat segment $\rightarrow (f')=0; (f'')=0$

f	0	0	0	0	0
f'		0	0	0	0
f''		0	0	0	




• Step $\rightarrow (f'): \{0, +, 0\}; (f''): \{0, +, -, 0\}$

f	0	0	0	7	7	7	7
f'		0	0	7	0	0	0
f''		0	7	-7	0	0	0



• Ramp $\rightarrow (f') \approx \text{constant}; (f'')=0$

f	5	4	3	2	1	0	0
f'	0	-1	-1	-1	-1	-1	0
f''	-1	0	0	0	0	1	0



The Laplacian Filter

- 2D second-order derivative operator used for image sharpening.
- Highlights sharp intensity transitions and de-emphasizes regions with slow intensity changes.
- Produces grayish edge lines and discontinuities on a dark background.

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

$$\therefore \frac{\partial^2 f}{\partial x^2} = f(x+1, y) + f(x-1, y) - 2f(x, y)$$

$$\therefore \frac{\partial^2 f}{\partial y^2} = f(x, y+1) + f(x, y-1) - 2f(x, y)$$

$$\nabla^2 f(x, y) = f(x+1, y) + f(x-1, y) + f(x, y+1) + f(x, y-1) - 4f(x, y)$$

- We can apply these kernels using convolution operations.

	$f(x-1, y)$	
$f(x, y-1)$	$f(x, y)$	$f(x, y+1)$
	$f(x+1, y)$	

Rotation invariant at 90° ++

0	1	0
1	-4	1
0	1	0

0	-1	0
-1	4	-1
0	-1	0

Laplacian kernel, Equation

Rotation invariant at 45° ++

1	1	1
1	-8	1
1	1	1

-1	-1	-1
-1	8	-1
-1	-1	-1

Includes the diagonal terms

Laplacian for Image Enhancement

To obtain the enhance image

$$g(x, y) = \begin{cases} f(x, y) - \nabla^2 f(x, y), & w_5 < 0 \\ f(x, y) + \nabla^2 f(x, y), & w_5 > 0 \end{cases}$$

- $f(x, y)$ → Original Image
- $\nabla^2 f(x, y)$ → Laplacian of original image
- In this way, background tonality can be perfectly preserved while details are enhanced.

Laplacian Kernel,
Highlight areas of
sharp intensity

-1	-1	-1
-1	8	-1
-1	-1	-1

+

0	0	0
0	1	0
0	0	0

Identity kernel, doesn't
modify the pixel values

w ₁	w ₂	w ₃
w ₄	w ₅	w ₆
w ₇	w ₈	w ₉

Resultant Kernel, produce a
sharper image by preserving
the original intensity values

-1	-1	-1
-1	9	-1
-1	-1	-1

Laplacian for Image Enhancement (Example)

$$g(x, y) = \begin{cases} f(x, y) - \nabla^2 f(x, y), & w_5 < 0 \\ f(x, y) + \nabla^2 f(x, y), & w_5 > 0 \end{cases}$$



Blurred image of the
North Pole of the
moon.



Laplacian image
obtained using the 90°
isotropic kernel



Image sharpened
using equation above



Image sharpened
using the same
procedure, but with
45° isotropic kernel.

The Gradient Filter

- 2D first derivative (∇f) in image processing are implemented using the magnitude of the gradient
- The gradient is generally given by

$$\nabla f = \begin{bmatrix} G_x \\ G_y \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

- The magnitude is given by

$$M(x, y) = \text{mag}(\nabla f) = [G_x^2 + G_y^2]^{\frac{1}{2}} = \left[\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2 \right]^{\frac{1}{2}}$$

Magnitude is approximated as

$$\nabla f \approx |G_x| + |G_y|$$

z_1	z_2	z_3
z_4	z_5	z_6
z_6	z_7	z_8

For a sub-image given in the figure:
The gradient is approximated by

$$\nabla f \approx |z_9 - z_5| + |z_8 - z_6|$$

-1	0
0	1

0	-1
1	0

This equation can be represented in masks (**Robert cross gradient operators**)

Cont.

Prewitt Operator

Normally the smallest mask used is of size 3 x 3

Based on the concept of approximating the gradient
several spatial masks have been proposed:

-1	-1	-1
0	0	0
1	1	1

Extract horizontal edges

-1	0	1
-1	0	1
-1	0	1

Extract vertical edges

z_1	z_2	z_3
z_4	z_5	z_6
z_6	z_7	z_8

Pixels arrangements

Prewitt Operator's Equation

$$\nabla f \approx |(z_7 + z_8 + z_9) - (z_1 + z_2 + z_3)| + |(z_3 + z_6 + z_9) - (z_1 + z_4 + z_7)|$$

Cont.

Sobel Operator

-1	-2	-1
0	0	0
1	2	1

Extract horizontal edges

-1	0	1
-2	0	2
-1	0	1

Extract vertical edges

z_1	z_2	z_3
z_4	z_5	z_6
z_6	z_7	z_8

Pixels arrangements


Prewitt Operator's Equation

$$\nabla f \approx |(z_7 + 2z_8 + z_9) - (z_1 + 2z_2 + z_3)| + |(z_3 + 2z_6 + z_9) - (z_1 + 2z_4 + z_7)|$$


↑
Emphasize more the current position (y)

↑
Emphasize more the current position (x)

Gradient Processing (example)




Optical image of contact lens
(note defects at 4 and 5 o'
clock)



Sobel Gradient

Sobel vs Laplacian



Thank You
For Your Attention